

# Exploratory\_data\_for\_GitHub\_Dataset

2024 年 3 月 26 日

## 1 Exploratory data analysis and data preprocessing for GitHub Dataset

### 1.1 一、数据摘要

```
[1]: # 首先导入相关包，观察数据结构
import pandas as pd
import numpy as np

target_data = pd.read_csv('./github_dataset.csv')
target_data.head()
```

```
[1]:
```

	repositories	stars_count	forks_count	issues_count	\
0	octocat/Hello-World	0	0	612	
1	EddieHubCommunity/support	271	150	536	
2	ethereum/aleth	0	0	313	
3	localstack/localstack	0	0	290	
4	education/classroom	0	589	202	

	pull_requests	contributors	language
0	316	2	NaN
1	6	71	NaN
2	27	154	C++
3	30	434	Python
4	22	67	Ruby

```
[2]: # 单独提出数据的属性列便于观察
target_data_columns = target_data.columns
target_data_columns
```

```
[2]: Index(['repositories', 'stars_count', 'forks_count', 'issues_count',
          'pull_requests', 'contributors', 'language'],
          dtype='object')
```

可以看到数据集已被成功加载，其属性列名称分别代表：

**repositories**（仓库名称），**stars\_count**（星标数），

**forks\_count**（分支数），**issues\_count**（问题数），

**pull\_requests**（拉取请求数），**contributors**（贡献者数），

以及 **language**（编程语言），从初步观察来看，**language** 列存在缺失值。

### 1. 标称属性，给出每个可能取值的频数

可以看到，标称属性只有 **repositories** 和 **language** 这两个，下面统计这两个标称属性的每个可能取值的频数

```
[3]: # 数据摘要：标称属性频数统计
nominal_columns = ['repositories', 'language']
for column in nominal_columns:
    print(f'Frequency of {column}:')
    print(target_data[column].value_counts(dropna=False))
    print('\n')
```

Frequency of repositories:

kameshsampath/ansible-role-rosa-demos	2
aloisdeniel/bluff	2
antoniaandreou/github-slideshow	2
jgthms/bulma-start	2
artkirienko/hlds-docker-dproto	2
..	..
WhiteHouse/CI0management	1
OxCaso/defillama-telegram-bot	1
ethereum/blake2b-py	1
openfoodfacts/folksonomy_mobile_experiment	1
gamemann/All_PropHealth	1
Name: repositories, Length: 972, dtype: int64	

Frequency of language:

JavaScript	253
------------	-----

Python	155
NaN	145
HTML	72
Java	44
CSS	37
TypeScript	37
Dart	36
Jupyter Notebook	29
C++	29
Ruby	28
C	26
Shell	25
PHP	16
Go	15
Swift	10
Rust	10
C#	8
Objective-C	8
Kotlin	7
Makefile	6
Jinja	5
SCSS	4
AutoHotkey	3
Dockerfile	3
CoffeeScript	3
Perl	3
Solidity	3
Vim Script	2
Pawn	2
Assembly	2
PowerShell	2
Hack	2
CodeQL	2
Vue	2
Elixir	2
Gherkin	1
QMake	1

CMake	1
Oz	1
Cuda	1
QML	1
ActionScript	1
Roff	1
HCL	1
R	1
PureBasic	1
Smarty	1
Less	1
Svelte	1
Haskell	1
SourcePawn	1

Name: language, dtype: int64

### 1.1.1 2. 数值属性，给出 5 数概括及缺失值的个数

除去上面两个标称属性外

对剩下的 stars\_count、forks\_count、issues\_count、pull\_requests、contributors 这五个数值属性进行五数概括，并计算缺失值的个数。

```
[4]: # 数值属性五数概括及缺失值个数
numerical_columns = ['stars_count', 'forks_count', 'issues_count', 'pull_requests', 'contributors']
print("Numerical columns summaries:")
for column in numerical_columns:
    print(f'{column}:')
    print(target_data[column].describe()) # 统计数值属性五数概括（最小值、三位四分位数、最大值）
    print(f'Missing values: {target_data[column].isnull().sum()}') # 统计数值属性缺失值个数
    print('\n')
```

Numerical columns summaries:

stars\_count:

```
count      1052.000000
mean        81.976236
std         170.403116
min          0.000000
25%          1.000000
50%         12.000000
75%         65.250000
max         995.000000
Name: stars_count, dtype: float64
Missing values: 0
```

```
forks_count:
count      1052.000000
mean        53.884981
std         127.699729
min          0.000000
25%          1.000000
50%          6.000000
75%         38.250000
max         973.000000
Name: forks_count, dtype: float64
Missing values: 0
```

```
issues_count:
count      1052.000000
mean         8.656844
std         32.445154
min          1.000000
25%          1.000000
50%          2.000000
75%          6.000000
max         612.000000
Name: issues_count, dtype: float64
Missing values: 0
```

```
pull_requests:
count      1052.000000
mean        4.374525
std         27.913732
min         0.000000
25%         0.000000
50%         0.000000
75%         2.000000
max         567.000000
Name: pull_requests, dtype: float64
Missing values: 0
```

```
contributors:
count      1052.000000
mean        8.364068
std         37.511807
min         0.000000
25%         0.000000
50%         2.000000
75%         4.000000
max         658.000000
Name: contributors, dtype: float64
Missing values: 0
```

经过上面的简要分析即可得知：

\* 标称属性 1. **repositories** 列有多个仓库名称，其中部分仓库名称出现频率超过 1 次，说明数据集中可能包含重复条目。2. **language** 列展示了不同的编程语言及其出现的频次，总计共 145 条记录的编程语言信息缺失。\* 数值属性 1. **stars\_count**：最小值 0，第一四分位数 1，中位数 12，第三四分位数 65.25，最大值 995，无缺失值。2. **forks\_count**：最小值 0，第一四分位数 1，中位数 6，第三四分位数 38.25，最大值 973，无缺失值。3. **issues\_count**：最小值 1，第一四分位数 1，中位数 2，第三四分位数 6，最大值 612，无缺失值。4. **pull\_requests**：最小值 0，第一四分位数 0，中位数 0，第三四分位数 2，最大值 567，无缺失值。5. **contributors**：最小值 0，第一四分位数 0，中位数 2，第三四分位数 4，最大值 658，无缺失值。

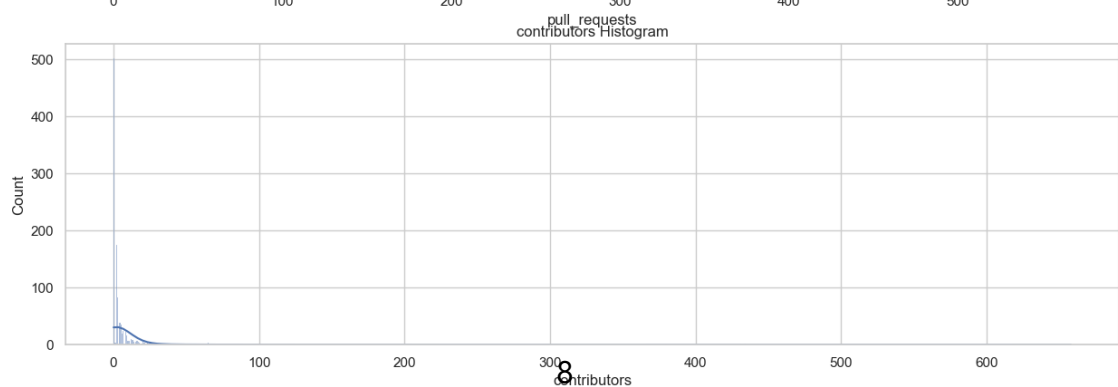




## 1.2 二、数据可视化

```
[5]: # 先导入相关包文件用于后续可视化
import matplotlib.pyplot as plt
import seaborn as sns

# 设置图形风格
sns.set(style="whitegrid")
```

### 1.2.1 1. 使用直方图检查数据分布

```
[6]: fig, axs = plt.subplots(5, 1, figsize=(15, 25))
# 绘制直方图
for i, column in enumerate(numerical_columns):
    sns.histplot(data=target_data, x=column, kde=True, ax=axs[i])
    axs[i].set_title(f'{column} Histogram')
```

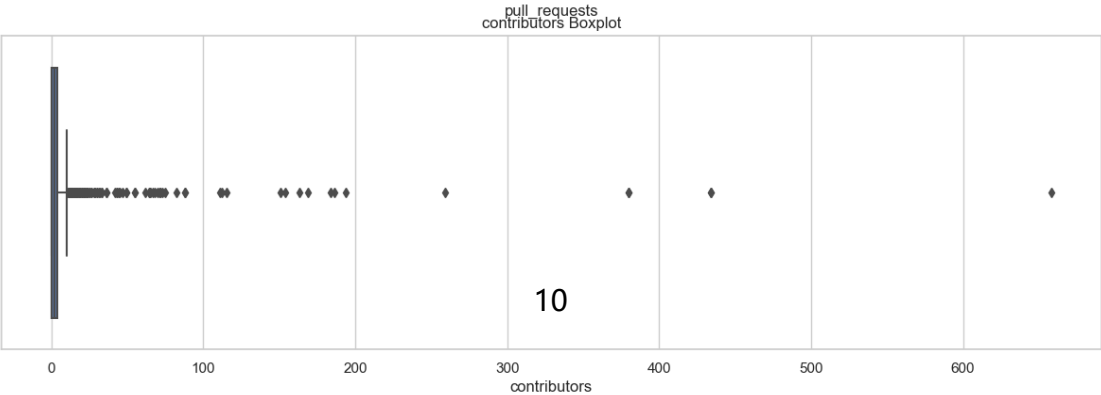
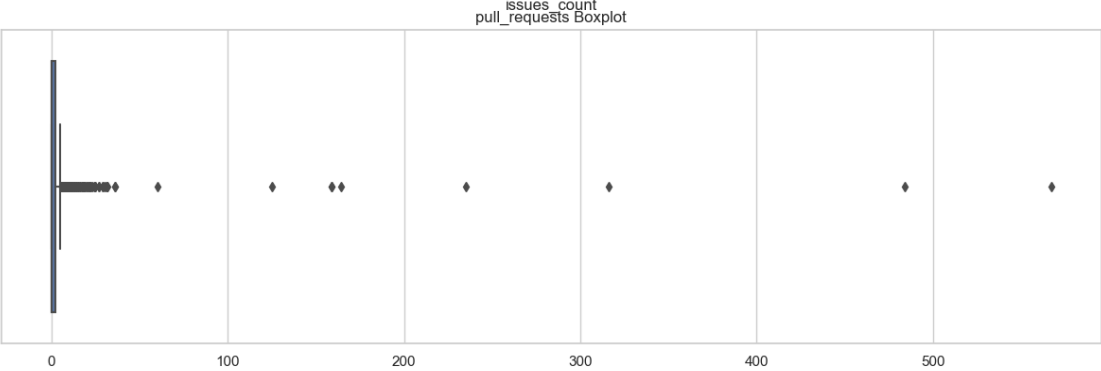
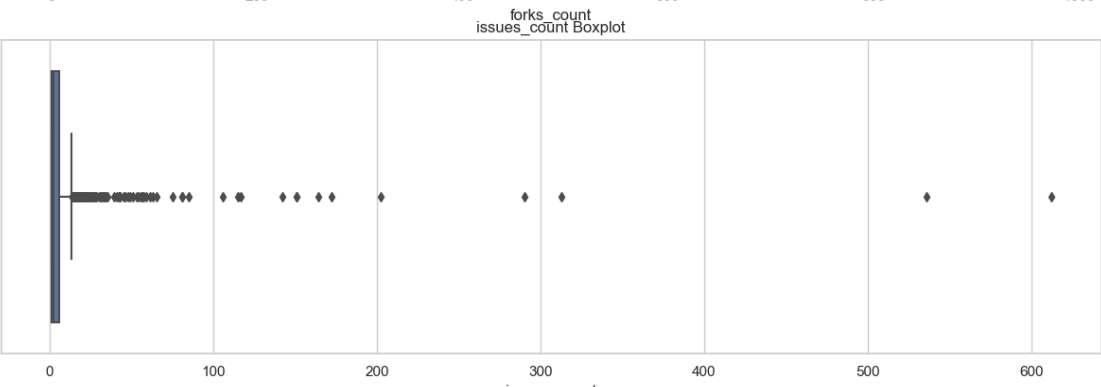
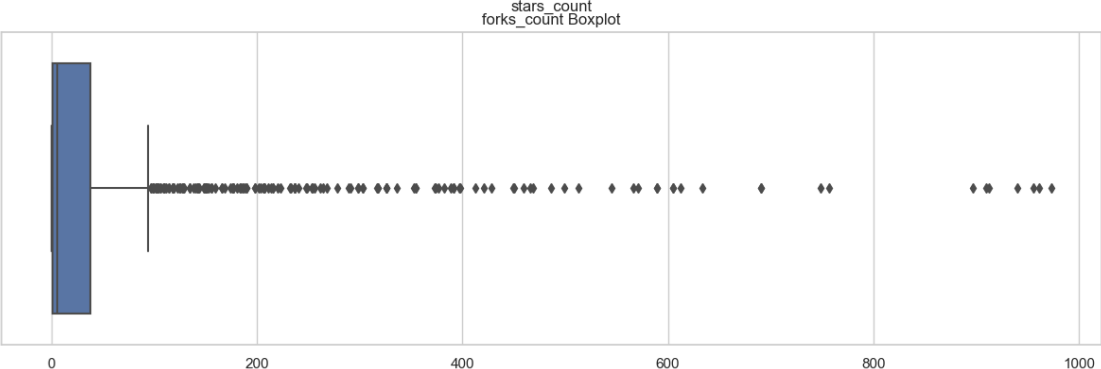
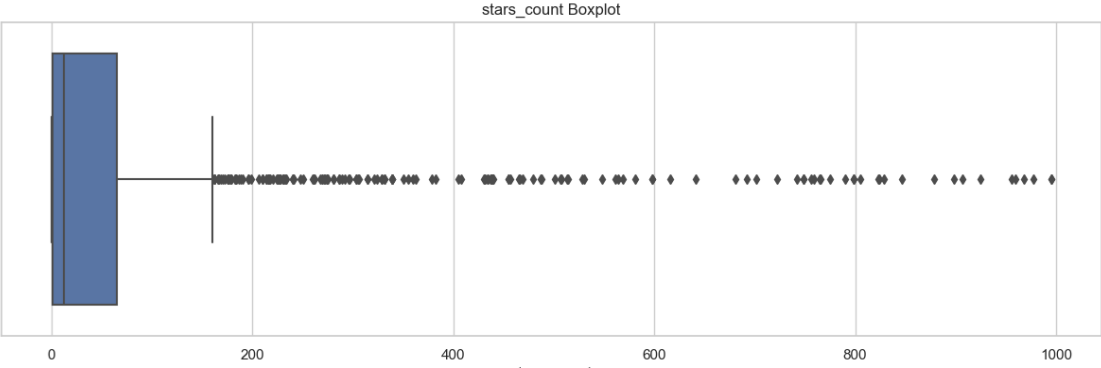




可以观察到，多数属性呈现右偏分布，  
意味着大多数仓库在星标数、分支数等指标上偏低，但存在少数高值的仓库。

### 1.2.2 2. 使用盒图检查数据离散点

```
[7]: fig, axs = plt.subplots(5, 1, figsize=(15, 25))
      # 绘制盒图
      for i, column in enumerate(numerical_columns):
          sns.boxplot(x=target_data[column], ax=axs[i])
          axs[i].set_title(f'{column} Boxplot')
```



可以观察到，几乎所有数值属性中都存在离散点，尤其是在 `stars_count`、`forks_count` 和 `contributors` 这几个指标上，离散点较为明显，这表明有些仓库在这些属性上远高于其他大多数仓库

## 2 三、缺失数据处理

### 2.0.1 1. 将缺失部分剔除

[8]: # 1. 剔除含有缺失值的记录

```
target_data_dropped = target_data.dropna()
target_data_dropped
```

```
[8]:
```

	repositories	stars_count	forks_count	\
2	ethereum/aleth	0	0	
3	localstack/localstack	0	0	
4	education/classroom	0	589	
5	shobhit97/open-gpstracker	0	0	
6	donnemartin/system-design-primer	0	0	
...	...	...	...	
1047	Tyriar/canvas-astar.dart	2	1	
1048	ankitkumar70777/github-slideshow	0	0	
1049	aitikgupta/interactive_cpu_scheduler	0	5	
1050	gwmccubbin/voting_dapp	11	5	
1051	gamemann/All_PropHealth	4	3	

	issues_count	pull_requests	contributors	language
2	313	27	154	C++
3	290	30	434	Python
4	202	22	67	Ruby
5	172	0	3	Java
6	164	164	115	Python
...	...	...	...	...
1047	1	0	0	Dart
1048	1	0	8	HTML
1049	1	1	7	Python

1050	1	0	0	JavaScript
1051	1	0	0	SourcePawn

[907 rows x 7 columns]

## 2.0.2 2. 用最高频率值来填补缺失值除

```
[9]: # 2. 使用最高频率值来填补缺失值
mode_language = target_data['language'].mode()[0] # 获取最高频率的语言
target_data_filled = target_data.fillna({'language': mode_language})

# 显示处理缺失值后的数据集大小变化
original_size = target_data.shape[0]
dropped_size = target_data_dropped.shape[0]
filled_size = target_data_filled.shape[0]
original_size, dropped_size, filled_size, mode_language
```

[9]: (1052, 907, 1052, 'JavaScript')

- 原始数据集大小：1052 条记录
- 剔除含有缺失值的记录后：数据集减少到 907 条记录
- 使用最高频率值填补缺失值后：数据集仍然保持 1052 条记录，但所有缺失的 language 值都被填补为最高频率的语言 'JavaScript'

```
[10]: target_data_filled
```

```
[10]:
```

	repositories	stars_count	forks_count	\
0	octocat/Hello-World	0	0	
1	EddieHubCommunity/support	271	150	
2	ethereum/aleth	0	0	
3	localstack/localstack	0	0	
4	education/classroom	0	589	
...	...	...	...	
1047	Tyriar/canvas-astar.dart	2	1	
1048	ankitkumar70777/github-slideshow	0	0	
1049	aitikgupta/interactive_cpu_scheduler	0	5	

1050	gwmccubbin/voting_dapp	11	5
1051	gamemann/All_PropHealth	4	3

	issues_count	pull_requests	contributors	language
0	612	316	2	JavaScript
1	536	6	71	JavaScript
2	313	27	154	C++
3	290	30	434	Python
4	202	22	67	Ruby
...	...	...	...	...
1047	1	0	0	Dart
1048	1	0	8	HTML
1049	1	1	7	Python
1050	1	0	0	JavaScript
1051	1	0	0	SourcePawn

[1052 rows x 7 columns]

### 2.0.3 3. 通过属性的相关关系来填补缺失值

```
[11]: target_data_coded = pd.get_dummies(target_data, columns=['language'], dummy_na=
↳ True, drop_first= True)
target_data_coded
```

```
[11]:
```

	repositories	stars_count	forks_count	\
0	octocat/Hello-World	0	0	
1	EddieHubCommunity/support	271	150	
2	ethereum/aleth	0	0	
3	localstack/localstack	0	0	
4	education/classroom	0	589	
...	...	...	...	
1047	Tyriar/canvas-astar.dart	2	1	
1048	ankitkumar70777/github-slideshow	0	0	
1049	aitikgupta/interactive_cpu_scheduler	0	5	
1050	gwmccubbin/voting_dapp	11	5	
1051	gamemann/All_PropHealth	4	3	

	issues_count	pull_requests	contributors	language_Assembly	\
0	612	316	2	0	
1	536	6	71	0	
2	313	27	154	0	
3	290	30	434	0	
4	202	22	67	0	
...	...	...	...	...	
1047	1	0	0	0	
1048	1	0	8	0	
1049	1	1	7	0	
1050	1	0	0	0	
1051	1	0	0	0	

	language_AutoHotkey	language_C	language_C#	...	language_Shell	\
0	0	0	0	...	0	
1	0	0	0	...	0	
2	0	0	0	...	0	
3	0	0	0	...	0	
4	0	0	0	...	0	
...	...	...	...	...	...	
1047	0	0	0	...	0	
1048	0	0	0	...	0	
1049	0	0	0	...	0	
1050	0	0	0	...	0	
1051	0	0	0	...	0	

	language_Smarty	language_Solidity	language_SourcePawn	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	...	...	...	
1047	0	0	0	
1048	0	0	0	
1049	0	0	0	

1050	0	0	0
1051	0	0	1

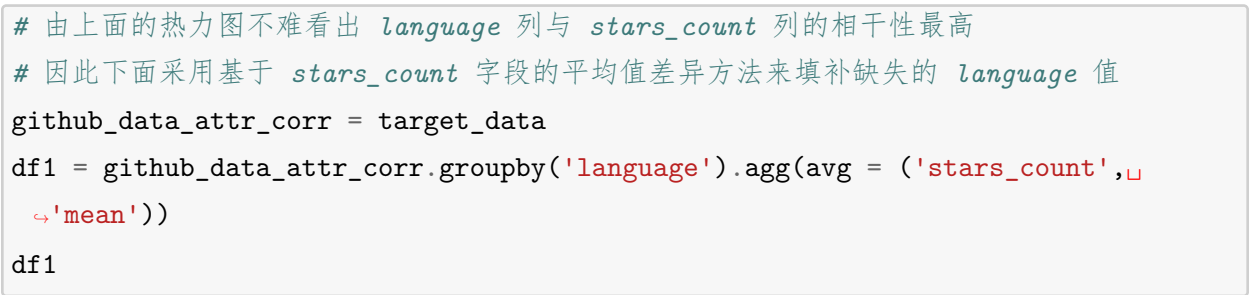
	language_Svelte	language_Swift	language_TypeScript	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	...	...	...	
1047	0	0	0	
1048	0	0	0	
1049	0	0	0	
1050	0	0	0	
1051	0	0	0	

	language_Vim Script	language_Vue	language_nan
0	0	0	1
1	0	0	1
2	0	0	0
3	0	0	0
4	0	0	0
...	...	...	...
1047	0	0	0
1048	0	0	0
1049	0	0	0
1050	0	0	0
1051	0	0	0

[1052 rows x 57 columns]

```
[12]: plt.figure(figsize=(40, 40))
sns.heatmap(target_data_coded.corr(method='spearman'), cmap='ocean', annot=True)
plt.title('Correlation Analysis')
```

```
[12]: Text(0.5, 1.0, 'Correlation Analysis')
```





[13]: avg

language	
ActionScript	3.000000
Assembly	426.000000
AutoHotkey	118.333333
C	153.076923
C#	43.250000
C++	39.172414
CMake	30.000000
CSS	49.675676
CodeQL	0.000000
CoffeeScript	7.333333
Cuda	14.000000
Dart	33.944444
Dockerfile	61.666667
Elixir	36.000000
Gherkin	1.000000
Go	139.600000
HCL	65.000000
HTML	54.888889
Hack	59.500000
Haskell	126.000000
Java	83.204545
JavaScript	87.083004
Jinja	0.400000
Jupyter Notebook	130.724138
Kotlin	26.714286
Less	0.000000
Makefile	0.833333
Objective-C	234.375000
Oz	37.000000
PHP	86.562500
Pawn	19.000000
Perl	294.000000
PowerShell	127.500000
PureBasic	5.000000

Python	73.038710
QML	24.000000
QMake	16.000000
R	1.000000
Roff	1.000000
Ruby	25.071429
Rust	56.400000
SCSS	228.500000
Shell	28.600000
Smarty	28.000000
Solidity	113.333333
SourcePawn	4.000000
Svelte	133.000000
Swift	152.500000
TypeScript	70.837838
Vim Script	602.000000
Vue	424.500000

```
[14]: from numpy import nan as NA
import warnings
warnings.filterwarnings(action='ignore')
for i in range(len(github_data_attr_corr)):
    if github_data_attr_corr['language'].iloc[i] is NA:
        rate = github_data_attr_corr['stars_count'].iloc[i]
        dist = []
        for j in range(len(df1)):
            dist.append(abs(df1.iloc[j]['avg']-rate))
        idx = dist.index(min(dist))
        github_data_attr_corr['language'].iloc[i] = df1.index[idx]
github_data_attr_corr['language'].value_counts()
```

```
[14]: JavaScript      257
Python              155
HTML                72
Java                44
CSS                 42
Dart                38
```

TypeScript	38
C	34
CodeQL	33
C++	33
Jupyter Notebook	29
Ruby	28
Shell	25
PHP	17
Go	17
Perl	14
ActionScript	12
Swift	10
Rust	10
C#	10
Kotlin	10
Vim Script	9
CoffeeScript	9
Pawn	8
Objective-C	8
Cuda	7
Gherkin	7
SCSS	7
Makefile	6
SourcePawn	6
AutoHotkey	5
QMake	5
Jinja	5
Assembly	4
PureBasic	4
Hack	4
QML	3
CMake	3
Dockerfile	3
HCL	3
Solidity	3
Elixir	2

```

Svelte                2
Smarty                2
Vue                   2
PowerShell             2
R                      1
Less                   1
Roff                   1
Haskell               1
Oz                     1
Name: language, dtype: int64

```

#### 2.0.4 4. 通过数据对象之间的相似性来填补缺失值 (拟使用热卡填充法, i.e. Hot Deck Imputation)

```

[15]: # 这里采用归一化巨鹿来作为相似性度量
      # 用这个相似性度量来找到最相似的数据对象

      # 选择用于计算相似性的数值属性
      attributes = ['stars_count', 'forks_count', 'issues_count', 'pull_requests', 'contributors']

      def regularit(df):
          new_df = pd.DataFrame(index=df.index)
          for c in attributes:
              d = df[c]
              MAX = d.max()
              MIN = d.min()
              new_df[c] = (d - MIN) / (MAX - MIN) # 归一化公式
          return new_df

```

```

[16]: github_data_sample_corr = target_data
      normal_github_data = regularit(github_data_sample_corr)
      normal_github_data

```

```

[16]:   stars_count  forks_count  issues_count  pull_requests  contributors
0      0.000000    0.000000    1.000000    0.557319    0.003040
1      0.272362    0.154162    0.875614    0.010582    0.107903
2      0.000000    0.000000    0.510638    0.047619    0.234043

```

3	0.000000	0.000000	0.472995	0.052910	0.659574
4	0.000000	0.605344	0.328969	0.038801	0.101824
...	...	...	...	...	...
1047	0.002010	0.001028	0.000000	0.000000	0.000000
1048	0.000000	0.000000	0.000000	0.000000	0.012158
1049	0.000000	0.005139	0.000000	0.001764	0.010638
1050	0.011055	0.005139	0.000000	0.000000	0.000000
1051	0.004020	0.003083	0.000000	0.000000	0.000000

[1052 rows x 5 columns]

```
[17]: normal_language = pd.concat([normal_github_data,
    ↪ github_data_sample_corr['language']], axis=1)
normal_language
```

```
[17]:
```

	stars_count	forks_count	issues_count	pull_requests	contributors	\
0	0.000000	0.000000	1.000000	0.557319	0.003040	
1	0.272362	0.154162	0.875614	0.010582	0.107903	
2	0.000000	0.000000	0.510638	0.047619	0.234043	
3	0.000000	0.000000	0.472995	0.052910	0.659574	
4	0.000000	0.605344	0.328969	0.038801	0.101824	
...	...	...	...	...	...	
1047	0.002010	0.001028	0.000000	0.000000	0.000000	
1048	0.000000	0.000000	0.000000	0.000000	0.012158	
1049	0.000000	0.005139	0.000000	0.001764	0.010638	
1050	0.011055	0.005139	0.000000	0.000000	0.000000	
1051	0.004020	0.003083	0.000000	0.000000	0.000000	

	language
0	CodeQL
1	Perl
2	C++
3	Python
4	Ruby
...	...
1047	Dart
1048	HTML

```

1049     Python
1050   JavaScript
1051   SourcePawn

```

```
[1052 rows x 6 columns]
```

```

[18]: infos = []
      for i in range(len(normal_language)):
          info = []
          star = normal_language['stars_count'].iloc[i]
          fork = normal_language['forks_count'].iloc[i]
          issue = normal_language['issues_count'].iloc[i]
          pull = normal_language['pull_requests'].iloc[i]
          contributor = normal_language['contributors'].iloc[i]
          info.append(star)
          info.append(fork)
          info.append(issue)
          info.append(pull)
          info.append(contributor)
          infos.append(info)

```

```

[19]: for i in range(len(normal_language)):
      if normal_language['language'].iloc[i] is NA:
          dists = []
          for j in len(rates):
              dist = np.sqrt(np.sum(np.square(infos[i] - infos[j])))
              dists.append(dist)
          idx = dists.index(min(dists))
          github_data_sample_corr['language'].iloc[i] =
      ↪github_data_sample_corr['language'].iloc[idx]
      github_data_sample_corr['language'].value_counts()

```

```

[19]: JavaScript      257
      Python          155
      HTML            72
      Java            44
      CSS             42

```

Dart	38
TypeScript	38
C	34
CodeQL	33
C++	33
Jupyter Notebook	29
Ruby	28
Shell	25
PHP	17
Go	17
Perl	14
ActionScript	12
Swift	10
Rust	10
C#	10
Kotlin	10
Vim Script	9
CoffeeScript	9
Pawn	8
Objective-C	8
Cuda	7
Gherkin	7
SCSS	7
Makefile	6
SourcePawn	6
AutoHotkey	5
QMake	5
Jinja	5
Assembly	4
PureBasic	4
Hack	4
QML	3
CMake	3
Dockerfile	3
HCL	3
Solidity	3

```
Elixir          2
Svelte          2
Smarty          2
Vue             2
PowerShell      2
R               1
Less            1
Roff            1
Haskell         1
Oz              1
Name: language, dtype: int64
```

### 3 四、预处理前后数据集差异比较分析

#### 3.0.1 1. 用最高频率值来填补缺失值除

```
[22]: language_counts_before = target_data['language'].value_counts().sort_index()
language_counts_after = target_data_filled['language'].value_counts().
      ↪sort_index()
```

```
[23]: # 合并前后填补数据，准备排序
combined_counts = pd.DataFrame({
    'Before Highest Frequency Filling': language_counts_before,
    'After Highest Frequency Filling': language_counts_after
})

# 按照填补前的频次进行降序排序
combined_counts_sorted = combined_counts.sort_values(by='Before Highest_
      ↪Frequency Filling', ascending=True)

# 绘制横向柱状图
plt.figure(figsize=(10, 15))
y_positions = range(len(combined_counts_sorted))
```




```

plt.barh(y_positions, combined_counts_sorted['Before Highest Frequency_
↳Filling'], height=0.4, label='Before Highest Frequency Filling',
↳color='blue', align='center')
plt.barh([p + 0.4 for p in y_positions], combined_counts_sorted['After Highest_
↳Frequency Filling'], height=0.4, label='After Highest Frequency Filling',
↳color='green', align='center')

plt.yticks([p + 0.2 for p in y_positions], combined_counts_sorted.index)
plt.xlabel('Frequency')
plt.ylabel('Language')
plt.legend()
plt.title('Language Frequency Before and After Highest Frequency Filling_
↳Missing Values')
plt.tight_layout()
plt.show()

```


# 可以看出这种简单粗暴的缺失值填补方式，拟合效果并不是很好



### 3.0.2 2. 通过属性的相关关系来填补缺失值

[24]: # 观察缺失值填补前后, *language* 频次的柱状图, 结果表明该方法可以很好的平滑原始分布。

```
language_counts = pd.DataFrame(target_data['language'].value_counts()).  
    ↳ sort_values(by='language', ascending=True).rename(columns={'language':  
    ↳ 'language_count'})  
language_count_attr_corr = language_counts  
language_count_attr_corr['language_count_attr_corr'] = [0] *  
    ↳ len(language_counts)  
  
for level in list(language_counts.index):  
    if level in list(github_data_attr_corr['language'].value_counts().index):  
        language_count_attr_corr.loc[[level], ['language_count_attr_corr']] =  
    ↳ github_data_attr_corr['language'].value_counts().loc[[level]].values[0]  
  
plt.figure(figsize=(40, 40))  
plt.yticks(fontsize=24)  
plt.xticks(fontsize=24)  
plt.barh(list(range(len(language_count_attr_corr))),  
    ↳ tick_label=language_count_attr_corr.index,  
    ↳ width=language_count_attr_corr['language_count'], label='language_count',  
    ↳ height=0.4, color='blue')  
plt.barh([d+0.42 for d in list(range(len(language_count_attr_corr)))],  
    ↳ tick_label=language_count_attr_corr.index,  
    ↳ width=language_count_attr_corr['language_count_attr_corr'],  
    ↳ label='language_count_attr_corr', height=0.4, color='red')  
plt.ylabel('language', fontsize=24)  
plt.xlabel('', fontsize=24)  
# plt.title('Number of movies for each appropriation-level?', fontsize=32,  
    ↳ loc='center')  
plt.legend(fontsize=32, loc='lower right')  
plt.show()
```



### 3.0.3 3. 通过数据对象之间的相似性来填补缺失值 (拟使用热卡填充法, i.e. Hot Deck Imputation)

```
[25]: language_counts_before_simulation = language_counts_before
language_counts_after_simulation = github_data_sample_corr['language'].
    ↪ value_counts()
# 将两个 Series 合并为 DataFrame 以便于排序和可视化
combined_counts_simulation = pd.DataFrame({
    'Before Similarity Filling': language_counts_before_simulation,
    'After Similarity Filling': language_counts_after_simulation
}).sort_values(by='Before Similarity Filling', ascending=True)
```

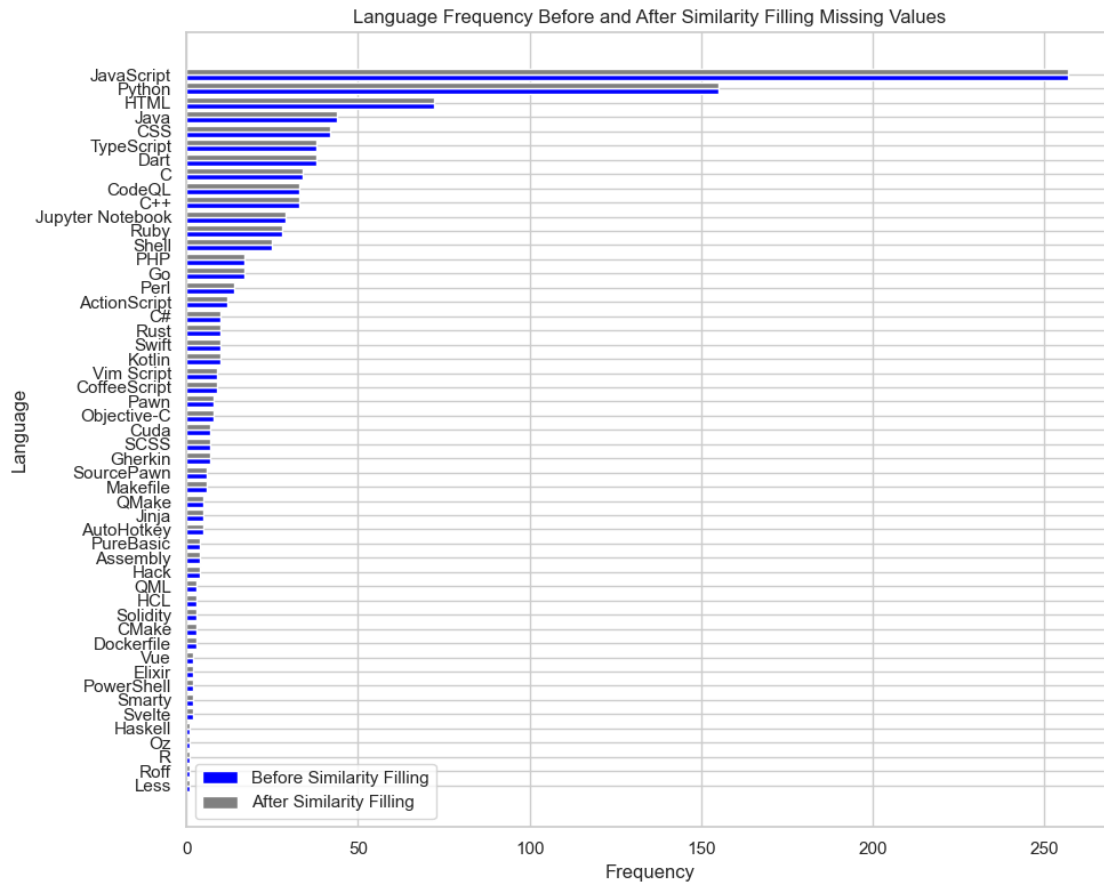
```

# 绘制横向柱状图
plt.figure(figsize=(10, 8))
y_positions = range(len(combined_counts_simulation))

plt.barh(y_positions, combined_counts_simulation['Before Similarity Filling'],
        height=0.4, label='Before Similarity Filling', color='blue', align='center')
plt.barh([p + 0.4 for p in y_positions], combined_counts_simulation['After_
        Similarity Filling'], height=0.4, label='After Similarity Filling',
        color='gray', align='center')

plt.yticks([p + 0.2 for p in y_positions], combined_counts_simulation.index)
plt.xlabel('Frequency')
plt.ylabel('Language')
plt.legend()
plt.title('Language Frequency Before and After Similarity Filling Missing_
        Values')
plt.tight_layout()
plt.show()
# 这种去除缺失值的方式也比较平滑可以看出

```



[ ]:

[ ]: